


Priority Queue Applications: Pathfinding in Weighted Graphs

EECS 214 Spring 2017

Priority Queues


- Like normal queues
 - objects wait in line to be processed
- But objects have an associated **numeric priority**
 - which is specified when the object is inserted
 - objects are dequeued in order of priority
- Different API
 - **Insert(object, priority)**
 - adds *object* with specified *priority*
 - **ExtractMax()**
 - except in most applications, we sort in ascending order, so...



http://www.alpcentauri.info/fig548_01_0.jpg


Priority Queues

- Like normal queues
 - objects wait in line to be processed
- But objects have an associated **numeric priority**
 - which is specified when the object is inserted
 - objects are dequeued in order of priority
- Different API
 - **Insert**(*object*, *priority*)
 - adds *object* with specified *priority*
 - **ExtractMin**()
 - **min** priority queues
 - returns **lowest** priority object




http://www.alpcentauri.info/fig548_01_0.jpg

Heaps Revisited



- Simple **tree structure** for implementing priority queues (they have other applications)
- We require that *parent nodes* **be larger than their children**
 - we don't require a sorted *in-order traversal*
- There are some crazyass types of heaps!
 - like the beap!
 - but we're going to focus on *binary heaps*
 - id est, **complete binary trees**
 - that also satisfy **the heap property**

Heaps Revisited





Proposition: the largest element of a heap is always its root

Proof:

- suppose some other element is the largest
- it *isn't* the root, so it has a parent
- it *is* the largest, so it is larger than its parent
- but this contradicts the definition of a heap, so this data structure is not a heap


Binary Heaps

- A **binary heap** is
 - a **complete binary tree**
 - that satisfies the **heap property**
-  LIT
 - so how do we represent one?




Embedding in an Array

- Turns out that a complete binary tree can be **embedded in an array** if we're clever about how we do it
- We can compute from an item's position
 - its parent's position
 - its children's positions



Embedding in an Array

- Store it in **breadth-first order**
 - store the **root** in the **first element** (element 0)
 - for any node:
 - let i be its array index
 - its **left child**: $2i + 1$
 - its **right child**: $2i + 2$
 - its **parent**: $(i - 1) / 2$



Insertion

HeapInsert(*A*, *Val*):

A.size++


i = *A*.size

while *i* > 0 and *A*[Parent(*i*)] < *Val*:

A[*i*] = *A*[Parent(*i*)]

i = Parent(*i*)

A[*i*] = *Val*



Extraction (Maximum)

HeapExtractMax(A):


max = A[0] // root === maximum

A[0] = A[A.size]

A.size--

Heapify(A, 0)

return max



Extraction (Maximum)

Heapify(A, i):

left = left(i); right = right(i)

if left <= A.size and A[left] > A[i]:

largest = left

else

largest = i


if right <= A.size and A[right] > A[largest]:

largest = r

if largest != i:


swap(A[i], A[largest])

Heapify(A, largest)



Finding the Shortest Path

- We've used **breadth-first search** to find a **shortest path** before
- But BFS assumes an equal weight for each edge
- When path planning, different edges get different **weights** or **costs**
- Costs can be:
 - time
 - distance
 - difficulty
 - anything, really, that you want to minimize



Pathfinding with Edge Costs

- BFS is super great, but it searches nodes naïvely
- So how do we make a version that searches nodes:
 - in order of **increasing total path cost**
 - rather than **increasing number of edges**
- That's right! **Priority queues!**
- A node's **priority** = **cost of path from start to that node**

Dijkstra's Least-cost Pathfinding Algorithm

Dijkstra(G, s, e):

PQ = new **PriorityQueue**() // starter bullshit

set all node **costs to infinity**

s.cost = 0

foreach node n in G:

 PQ.**Insert**(n, n.cost)

while not PQ.IsEmpty:

 u = PQ.**ExtractMin**()

 if u == e: return;

 foreach neighbor in **u.neighbors**:

 w = edge(u, v).weight // edge weight between nodes u and v

newcost = **u.cost** + w

 if newcost < v.cost:

 PQ.**DecreaseKey**(v, newcost)

 v.cost = newcost

 v.predecessor = u

Hold up. Decrease Key?

- We need a new operation in our API that lets us change our priority queue:
decreasing the priority of a given object already in the queue
 - N.B.: we're using a min-priority queue ('cause we're doing `ExtractMin`), so decreasing a queue moves it *forward* in the priority queue
- Now, how should we implement `DecreaseKey`?

Implementing DecreaseKey

- One option:
 - we could **remove it**
 - then **reinsert it at a lower priority**
- But the insertion algorithm:
 - adds at the bottom of the heap
 - then **swaps it upward** until its priority is lower than its parent's (for a min heap)

HeapInsert(A, Val):

A.size++

i = A.size

while i > 0 and A[Parent(i)] > Val:

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = Val

N.B: we're checking if *Parent* > *Val* now; it's a min-priority queue

Implementing DecreaseKey

- So DecreaseKey is actually super easy:
 - just move a node up
 - until it's in the right place
- We're going to copy the code for insert and remove the bits that put it at the end

DecreaseKey(A, i, Val):

while $i > 0$ and $A[\text{Parent}(i)] > Val$:

$A[i] = A[\text{Parent}(i)]$

$i = \text{Parent}(i)$

$A[i] = Val$

Implementing DecreaseKey

- Hey, we added an argument!
- Unfortunately, we need to remember **where the node** is in the damn heap
 - id est, its **array index**
- Best done by storing it in the **graph node itself**
- Have to remember to **update it** any time the node moves around in the queue

DecreaseKey(A, i, Val):

while $i > 0$ and $A[\text{Parent}(i)] > Val$:

$A[i] = A[\text{Parent}(i)]$ // here

$i = \text{Parent}(i)$

$A[i] = Val$ // and here

Proving Correctness

- There are two kinds of nodes:
 - those still **in the queue**
 - those that have been **removed from the queue**

Dijkstra(G, s, e):

PQ = new **PriorityQueue**()

set all node **costs to infinity**

s.cost = 0

foreach node n in G:

PQ.**Insert**(n, n.cost)

while not PQ.IsEmpty:

u = PQ.**ExtractMin**()

if u == e:

return;

foreach neighbor in **u.neighbors**:

w = edge(u, v).weight

newcost = u.cost + w

if newcost < v.cost:

PQ.**DecreaseKey**(v, newcost)

v.cost = newcost

v.predecessor = u

Invariants

For all nodes V :

- if V is **not in the queue**
 - $\text{dist}(V)$ is the length of the **shortest path** from the start node to V
- otherwise
 - $\text{dist}(V)$ is the length of the **shortest path** from the start node to V using **nodes not in the queue**

Dijkstra(G, s, e):

PQ = new **PriorityQueue**()

set all node **costs to infinity**

$s.\text{cost} = 0$

foreach node n in G :

PQ.**Insert**($n, n.\text{cost}$)

while not PQ.IsEmpty:

$u = \text{PQ}.\text{ExtractMin}()$

if $u == e$:

return;

foreach neighbor in **$u.\text{neighbors}$** :

$w = \text{edge}(u, v).\text{weight}$

$\text{newcost} = u.\text{cost} + w$

if $\text{newcost} < v.\text{cost}$:

PQ.**DecreaseKey**($v, \text{newcost}$)

$v.\text{cost} = \text{newcost}$

$v.\text{predecessor} = u$

Proof

For all nodes V :

- if V is **not in the queue**
 - $\text{dist}(V)$ is the length of the **shortest path** from the start node to V
- otherwise
 - $\text{dist}(V)$ is the length of the **shortest path** from the start node to V using **nodes not in the queue**

Base Case: **only the start node** has been removed from the queue

Trivially true:

- $\text{dist}(\text{start}) = 0$
- $\text{dist}(\text{anything else}) = \text{float.infinity}$

Proof

For all nodes V :

- if V is **not in the queue**
 - $\text{dist}(V)$ is the length of the **shortest path** from the start node to V
- otherwise
 - $\text{dist}(V)$ is the length of the **shortest path** from the start node to V using **nodes not in the queue**

Inductive case:

- assume it's **true for the currently visited nodes**
- let V be the **next node** in the queue

Then:

1. V has the **minimum $\text{dist}()$** of **all the nodes in the queue**
2. $\text{dist}(V)$ is the length of the shortest path to V using nodes not in the queue

Proof

Inductive case:

- assume it's true for the currently visited nodes
- let V be the next node in the queue

Then:

1. v has the minimum $\text{dist}()$ of all the nodes in the queue
2. $\text{dist}(V)$ is the length of the shortest path to V using nodes not in the queue

Claim:

- $\text{dist}(V)$ is the length of the shortest path using any nodes

Proof:

- Assume there exists a shorter path
- Case 1: the shorter path *contains a node in the queue*
 - then that node must be closer to the start than V
 - this contradicts (1)

Proof

Inductive case:

- assume it's true for the currently visited nodes
- let V be the next node in the queue

Then:

1. v has the minimum $\text{dist}()$ of all the nodes in the queue
2. $\text{dist}(V)$ is the length of the shortest path to V using nodes not in the queue

Claim:

- $\text{dist}(V)$ is the length of the shortest path using any nodes

Proof:

- Assume there exists a shorter path
- Case 2: the shorter path *contains no nodes from the queue*
 - this contradicts (2)

Proof

Inductive case:

- assume it's true for the currently visited nodes
- let V be the next node in the queue


Then:

1. v has the minimum $\text{dist}()$ of all the nodes in the queue
2. $\text{dist}(V)$ is the length of the shortest path to V using nodes not in the queue

Therefore:

- The inductive case holds
- The invariants hold
 - including when the queue is empty
 - at which point we know the distance to every node
- The algorithm is correct


Running Dijkstra's



Initialize Priority Queue

Queue:


1. A (0)
2. B (∞)
3. C (∞)
4. D (∞)
5. E (∞)
6. F (∞)



Extract Min (A)

Queue:

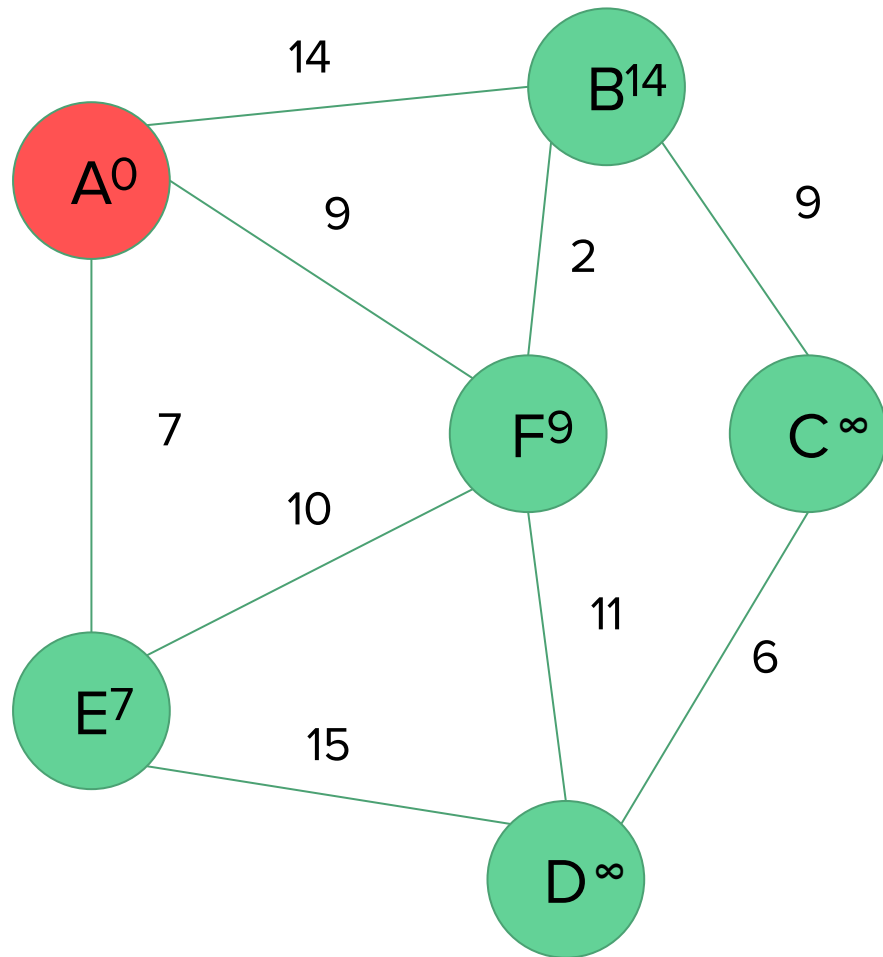
1. B (∞)
2. C (∞)
3. D (∞)
4. E (∞)
5. F (∞)



Update Neighbors

Queue:


1. B (14)
2. C (∞)
3. D (∞)
4. E (7)
5. F (9)



Update Neighbors

Queue:

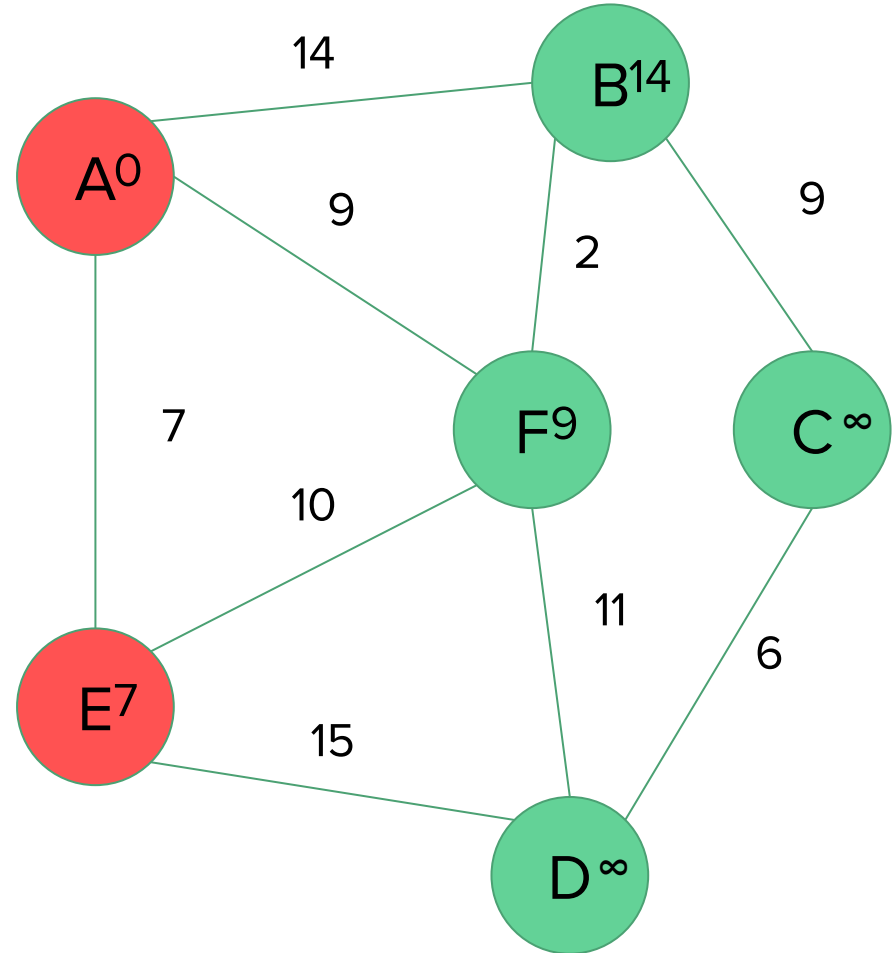
1. E (7)
2. F (9)
3. B (14)
4. C (∞)
5. D (∞)



Extract Min (E)

Queue:


1. F (9)
2. B (14)
3. C (∞)
4. D (∞)



Update Neighbors

Queue:


1. F (9)
2. B (14)
3. C (∞)
4. D (22)



Update Neighbors

Queue:


1. F (9)
2. B (14)
3. D (22)
4. C (∞)



Extract Min (F)

Queue:


1. B (14)
2. D (22)
3. C (∞)



Update Neighbors

Queue:


1. B (11)
2. D (20)
3. C (∞)



Extract Min (B)

Queue:

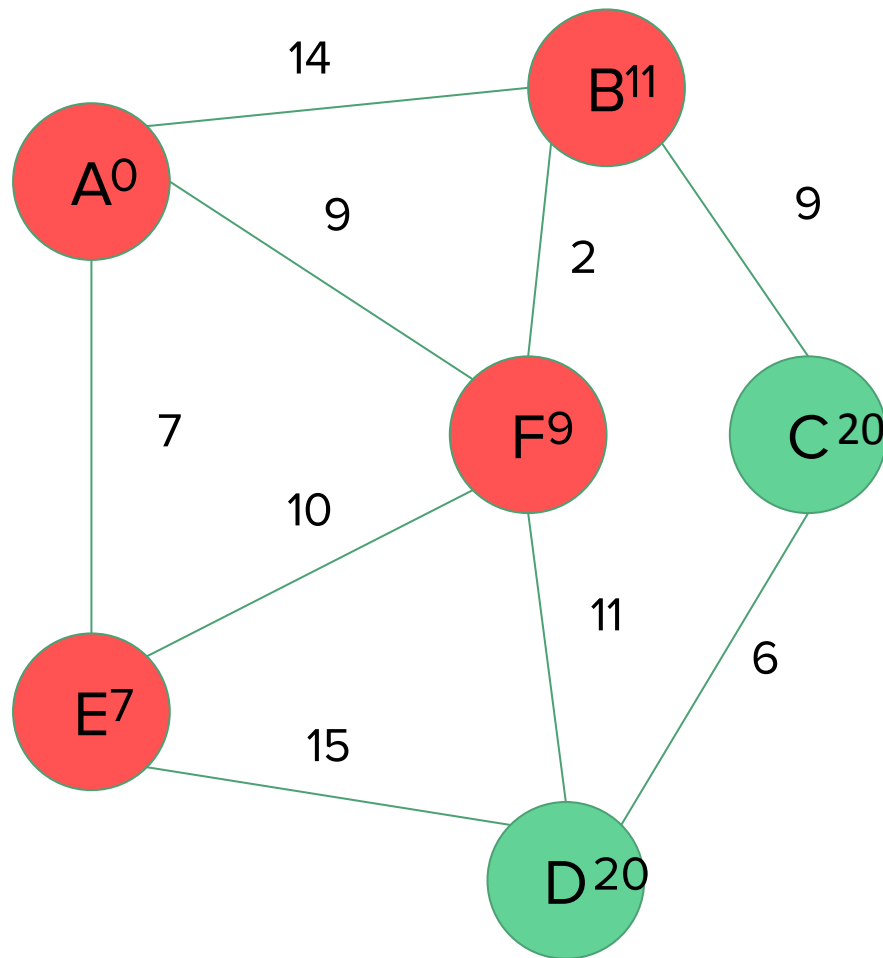
1. D (20)
2. C (∞)



Update Neighbors


Queue:

1. D (20)
2. C (20)




We Reached the Goal

So we can empty our queue, since we know that we have the shortest path.



Shortest Path



Runtime

Dijkstra(G, s, e):

PQ = new **PriorityQueue**() _____

set all node **costs to infinity** _____

$s.cost = 0$ _____

foreach node n in G :

 PQ.**Insert**($n, n.cost$) _____

while not PQ.IsEmpty:

$u = \text{PQ.ExtractMin}()$ _____

 if $u == e$: return;

 foreach neighbor in **$u.neighbors$** :

$w = \text{edge}(u, v).weight$ _____

$newcost = u.cost + w$

 if $newcost < v.cost$:

 PQ.**DecreaseKey**($v, newcost$)

$v.cost = newcost$

$v.predecessor = u$

Runs:

Once

Once

Once

$O(V)$ times

$O(V)$ times

$O(V)$

$O(E)$ times*

$O(E)$

$O(E)$

$O(E)$

$O(E)$

Runtime

Dijkstra(G, s, e):

PQ = new **PriorityQueue**() _____

set all node **costs to infinity** _____

$s.cost = 0$ _____

foreach node n in G :

 PQ.**Insert**($n, n.cost$) _____

while not PQ.IsEmpty:

$u = \text{PQ.ExtractMin}()$ _____

 if $u == e$: return;

 foreach neighbor in **$u.neighbors$** :

$w = \text{edge}(u, v).weight$ _____

$newcost = u.cost + w$

 if $newcost < v.cost$:

 PQ.**DecreaseKey**($v, newcost$)

$v.cost = newcost$

$v.predecessor = u$

Runs:

Once

Once

Once

$O(V)$ times

$O(V)$ times

$O(V)$

$O(E)$ times*

$O(E)$

$O(E)$

$O(E)$

$O(E)$

*Why not $O(EV)$?

Runtime

Dijkstra(G, s, e):

PQ = new **PriorityQueue**() _____

set all node **costs to infinity** _____

s.cost = 0 _____

foreach node n in G:

 PQ.**Insert**(n, n.cost) _____

while not PQ.IsEmpty:

 u = PQ.**ExtractMin**() _____

 if u == e: return;

 foreach neighbor in **u.neighbors**:

 w = edge(u, v).weight _____

newcost = **u.cost** + w

 if newcost < v.cost:

 PQ.**DecreaseKey**(v, newcost)

 v.cost = newcost

 v.predecessor = u

Runs:

Once

Once

Once

$O(V)$ times

$O(V)$ times

$O(V)$

$O(E)$ times

$O(E)$

$O(E)$

$O(E)$

$O(E)$

Execution Time:

$O(1)$

$O(V)$

$O(1)$

$O(\log V)$

$O(\log V)$

$O(1)$

$O(1)$

$O(1)$

$O(\log V)$

$O(1)$

$O(1)$

Runtime

Dijkstra(G, s, e):

PQ = new **PriorityQueue**() _____

set all node **costs to infinity** _____

s.cost = 0 _____

foreach node n in G:

 PQ.**Insert**(n, n.cost) _____

while not PQ.IsEmpty:

 u = PQ.**ExtractMin**() _____

 if u == e: return;

 foreach neighbor in **u.neighbors**:

 w = edge(u, v).weight _____

newcost = **u.cost** + w

 if newcost < v.cost:

 PQ.**DecreaseKey**(v, newcost)

 v.cost = newcost

 v.predecessor = u

Runs:

Once

Once

Once

$O(V)$ times

$O(V)$ times

$O(V)$

$O(E)$ times

$O(E)$

$O(E)$

$O(E)$

$O(E)$

Execution Time:

$O(1)$

$O(V)$

$O(1)$

$O(\log V)$

$O(\log V)$

$O(1)$

$O(1)$

$O(1)$

$O(\log V)$

$O(1)$

$O(1)$

Total:

$O(1)$

$O(V)$

$O(1)$

$O(V \log V)$

$O(V \log V)$

$O(V)$

$O(E)$

$O(E)$

$O(E \log V)$

$O(E)$


$O(E)$

Runtime

$$\begin{aligned} &O(1) + O(V) + O(E) + O(V \log V) + O(E \log V) \\ &= O(1) + O(V + E) + O((E + V) \log V) \\ &= O(V + E) + O((E + V) \log V) \\ &= O((V + E) (1 + \log V)) \\ &= O((V + E) \log V) \end{aligned}$$

N.B.: you can *technically* get better runtime $[O(E + V \log V)]$ if you use something called a Fibonacci heap, but it's so laughably theoretical that this might as well be the best that we've got.

Example(ish): Searching in a 4-Connected Grid



Source:
<http://theory.stanford.edu/~amitp/GameProgramming/AStarComparison.html>

Dijkstra's Does a lot of Useless Searching

- It's a **blind search** algorithm
 - it **exhaustively** explores nodes in **increasing order of distance**
 - even if they're in the opposite direction of the goal! D:
- Is there some way of **biasing** the search towards more productive nodes?
 - search the **more promising** nodes first




Heuristic Search

- Dijkstra's sorts the queue by distance from the start node
 - and so searches in order of increasing distance from the start
- Suppose you had some estimate $h(v)$ of v 's distance from the goal
 - we call that a heuristic
- Now we sort the priority queue by estimated total distance
 - (known) distance from start + estimated distance to the goal
 - id est: $\text{dist}(V) + h(V)$

A*

- Almost identical to Dijkstra's, but uses $\text{dist}(V) + h(V)$ to sort the priority queue
- Guaranteed to find the optimal path provided that $h(V)$ never overestimates the distance of V to the goal
 - referred to as an optimistic algorithm
 - the proof is basically the same as for Dijkstra's
 - in fact, you may notice that Dijkstra's algorithm is just A* where the heuristic functions is constant
- The standard algorithm for computing paths in space
 - e.g., Google Maps, robotics, game AI, etc.
 - there are others, but this is bread and butter
- Most common heuristic is $h(V) = \| \text{position}(V) - \text{position}(\text{goal}) \|$

A^* is Way More Efficient (if you've got a heuristic)



Reading

- I can't (and don't want to) assign y'all homework, but this was in Ian's slide deck so here is a relevant reading that I assume that he wants you to read:
 - CLR chapter 24.3 (Dijkstra's Algorithm)